

Practicum_1

Minxin Cheng

0. Load packages

```
#install.packages("DMwR")  
  
library(dplyr)  
library(ggplot2)  
library(ggpubr)  
library(reshape2)  
require(zoo)  
library(class)  
library(gmodels)  
library(DMwR)
```

Problem 1

Question 1

Download the data set (glass data) along with its explanation. Note that the data file does not contain header names, you may wish to add those. The description of each column can be found in the data set explanation.

```
# read csv  
glassData <- read.csv("glass.data",  
                      header = FALSE)  
  
# assign column names  
colnames(glassData) <- c("id", "refIdx", "Na", "Mg",  
                        "Al", "Si", "K", "Ca", "Ba",  
                        "Fe", "type")
```

Question 2

Explore the data set to get a sense of the data and to get comfortable with it.

```
# check if there is any missing values  
any(is.na(glassData))
```

```
## [1] FALSE
```

```
# check if the data is structured  
str(glassData)
```

```
## 'data.frame': 214 obs. of 11 variables:
## $ id : int 1 2 3 4 5 6 7 8 9 10 ...
## $ refIdx: num 1.52 1.52 1.52 1.52 1.52 ...
## $ Na : num 13.6 13.9 13.5 13.2 13.3 ...
## $ Mg : num 4.49 3.6 3.55 3.69 3.62 3.61 3.6 3.61 3.58 3.6 ...
## $ Al : num 1.1 1.36 1.54 1.29 1.24 1.62 1.14 1.05 1.37 1.36 ...
## $ Si : num 71.8 72.7 73 72.6 73.1 ...
## $ K : num 0.06 0.48 0.39 0.57 0.55 0.64 0.58 0.57 0.56 0.57 ...
## $ Ca : num 8.75 7.83 7.78 8.22 8.07 8.07 8.17 8.24 8.3 8.4 ...
## $ Ba : num 0 0 0 0 0 0 0 0 0 0 ...
## $ Fe : num 0 0 0 0 0 0.26 0 0 0 0.11 ...
## $ type : int 1 1 1 1 1 1 1 1 1 1 ...
```

```
# check each column's type see if anyone needs to be changed
sapply(glassData, class)
```

```
##      id      refIdx      Na      Mg      Al      Si      K      Ca
## "integer" "numeric" "numeric" "numeric" "numeric" "numeric" "numeric" "numeric"
##      Ba      Fe      type
## "numeric" "numeric" "integer"
```

```
# get an overall information of the data set
summary(glassData)
```

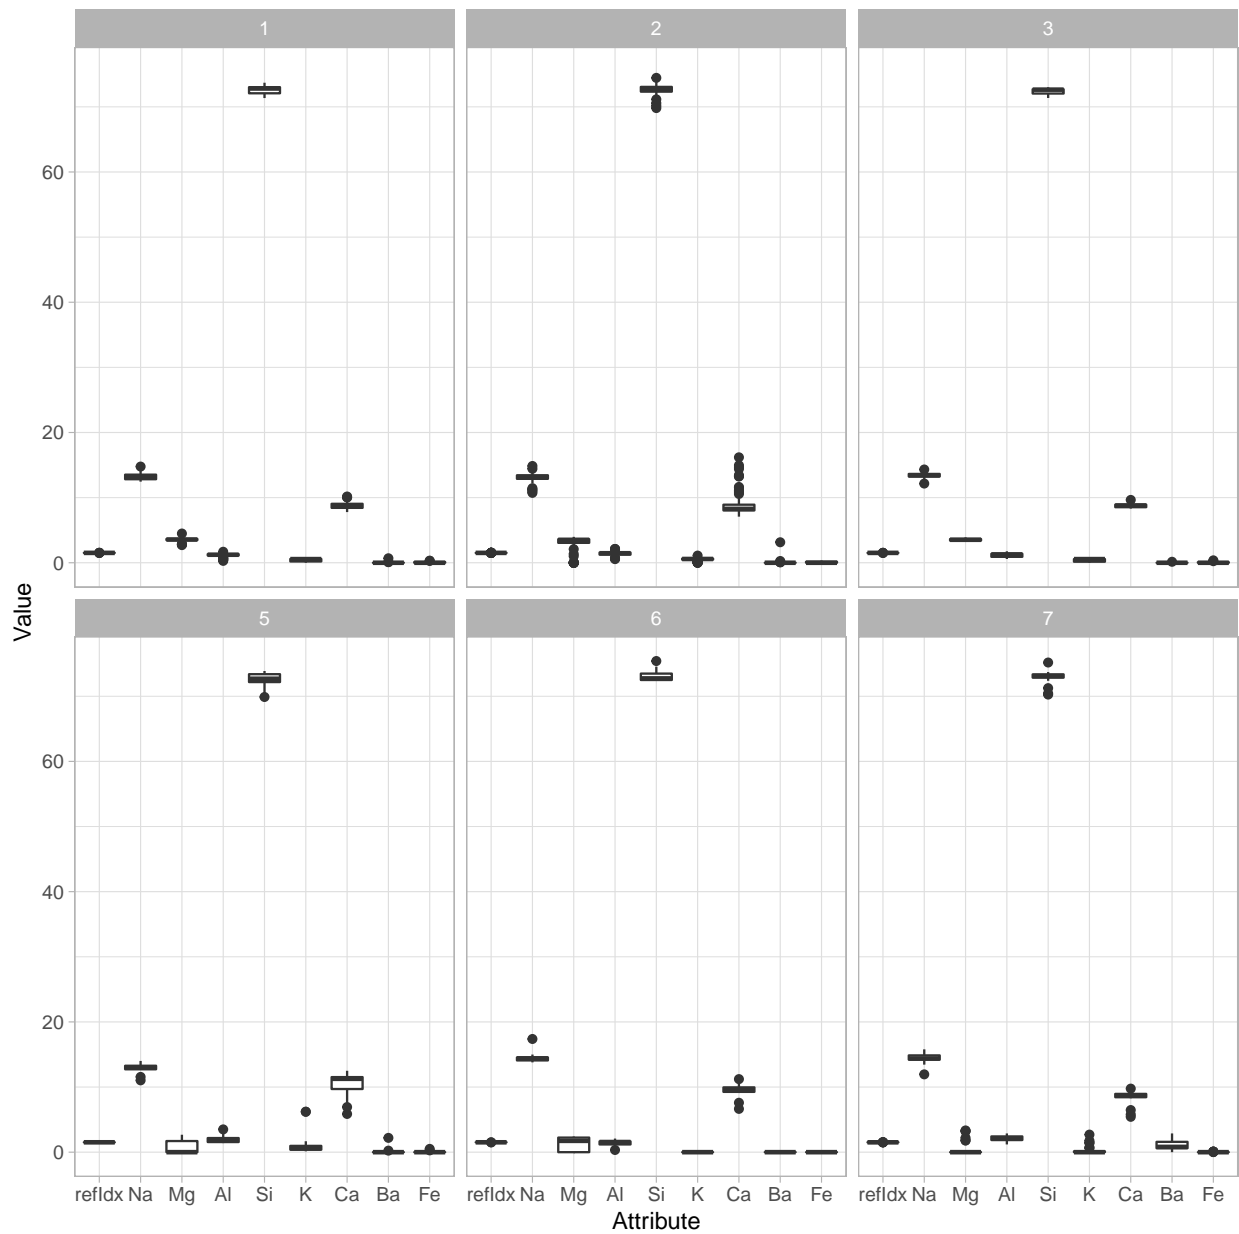
```
##      id      refIdx      Na      Mg
## Min.   : 1.00   Min.   :1.511   Min.   :10.73   Min.   :0.000
## 1st Qu.: 54.25   1st Qu.:1.517   1st Qu.:12.91   1st Qu.:2.115
## Median :107.50   Median :1.518   Median :13.30   Median :3.480
## Mean   :107.50   Mean   :1.518   Mean   :13.41   Mean   :2.685
## 3rd Qu.:160.75   3rd Qu.:1.519   3rd Qu.:13.82   3rd Qu.:3.600
## Max.   :214.00   Max.   :1.534   Max.   :17.38   Max.   :4.490
##      Al      Si      K      Ca
## Min.   :0.290   Min.   :69.81   Min.   :0.0000   Min.   : 5.430
## 1st Qu.:1.190   1st Qu.:72.28   1st Qu.:0.1225   1st Qu.: 8.240
## Median :1.360   Median :72.79   Median :0.5550   Median : 8.600
## Mean   :1.445   Mean   :72.65   Mean   :0.4971   Mean   : 8.957
## 3rd Qu.:1.630   3rd Qu.:73.09   3rd Qu.:0.6100   3rd Qu.: 9.172
## Max.   :3.500   Max.   :75.41   Max.   :6.2100   Max.   :16.190
##      Ba      Fe      type
## Min.   :0.000   Min.   :0.00000   Min.   :1.00
## 1st Qu.:0.000   1st Qu.:0.00000   1st Qu.:1.00
## Median :0.000   Median :0.00000   Median :2.00
## Mean   :0.175   Mean   :0.05701   Mean   :2.78
## 3rd Qu.:0.000   3rd Qu.:0.10000   3rd Qu.:3.00
## Max.   :3.150   Max.   :0.51000   Max.   :7.00
```

```
# box plots to get an overall sense by glass type
## extract all glass type information
glassType <- glassData %>%
  select(id, type)
## reshape the data
glassOverview <- melt(glassData[, 1:10],
```

```

        id.vars = "id",
        variable.name = "attribute")
## add type information to the reshaped data
glassOverview <- glassOverview %>%
  left_join(glassType, by = "id")
## box plot
ggplot(glassOverview, aes(x = attribute, y = value)) +
  geom_boxplot() +
  facet_wrap(~type) +
  labs(x = "Attribute", y = "Value") +
  theme_light()


```



Question 3

Create a histogram of column 2 (refractive index) and overlay a normal curve.

```
ggplot(glassData, aes(x = refIdx)) +  
  geom_histogram(breaks = seq(1.510, 1.540, by = 0.001),  
    colour = "black",  
    fill = "grey") +  
  stat_function(fun = dnorm,  
    args = list(mean = mean(glassData$refIdx),  
      sd = sd(glassData$refIdx))) +  
  theme_light()
```




Question 4 Test normality of column 2 by performing either a Shapiro-Wilk or Kolmogorov-Smirnov test. Describe what you found.

```
# run shapiro test  
shapiro.test(glassData$refIdx)
```

```
##  
## Shapiro-Wilk normality test  
##  
## data:  glassData$refIdx  
## W = 0.86757, p-value = 1.077e-12
```

```
# do a qq plot to visualize
ggqqplot(glassData$refIdx)
```



As the Shapiro test result, $p < 0.05$, therefore column 2 is not normally distributed. Q-Q plot also supported it.

Question 5

Identify any outliers for the columns using a z-score deviation approach. i.e., consider any values that are more than 2 standard deviations from the mean as outliers. Which are your outliers for each column? What would you do? Summarize potential strategies in your notebook.

```
# normalize all columns of the dataset except id and type columns
glassDataNorm <- as.data.frame(scale(glassData[, 2:10]))
# summary the normalized dataset to get an overview
summary(glassDataNorm)
```

##	refIdx	Na	Mg	Al
##	Min. : -2.3759	Min. : -3.2793	Min. : -1.8611	Min. : -2.3132
##	1st Qu.: -0.6069	1st Qu.: -0.6127	1st Qu.: -0.3948	1st Qu.: -0.5106
##	Median : -0.2257	Median : -0.1321	Median : 0.5515	Median : -0.1701
##	Mean : 0.0000	Mean : 0.0000	Mean : 0.0000	Mean : 0.0000
##	3rd Qu.: 0.2608	3rd Qu.: 0.5108	3rd Qu.: 0.6347	3rd Qu.: 0.3707
##	Max. : 5.1252	Max. : 4.8642	Max. : 1.2517	Max. : 4.1162
##	Si	K	Ca	Ba

```
## Min.      :-3.6679   Min.      :-0.76213   Min.      :-2.4783   Min.      :-0.3521
## 1st Qu.: -0.4789   1st Qu.: -0.57430   1st Qu.: -0.5038   1st Qu.: -0.3521
## Median : 0.1795   Median : 0.08884   Median : -0.2508   Median : -0.3521
## Mean    : 0.0000   Mean    : 0.00000   Mean    : 0.0000   Mean    : 0.0000
## 3rd Qu.: 0.5636   3rd Qu.: 0.17318   3rd Qu.: 0.1515   3rd Qu.: -0.3521
## Max.     : 3.5622   Max.     : 8.75961   Max.     : 5.0824   Max.     : 5.9832
##          Fe
## Min.      :-0.5851
## 1st Qu.: -0.5851
## Median : -0.5851
## Mean     : 0.0000
## 3rd Qu.: 0.4412
## Max.     : 4.6490
```

```
sapply(glassDataNorm, class)
```


```
## refIdx      Na      Mg      Al      Si      K      Ca      Ba
## "numeric" "numeric" "numeric" "numeric" "numeric" "numeric" "numeric" "numeric"
##          Fe
## "numeric"
```

```
# find out outliers as long as there is one or more columns value is more than
# 2 standard deviations from the mean
outlier <- glassDataNorm %>%
  filter_all(any_vars(abs(.) > 2))
```

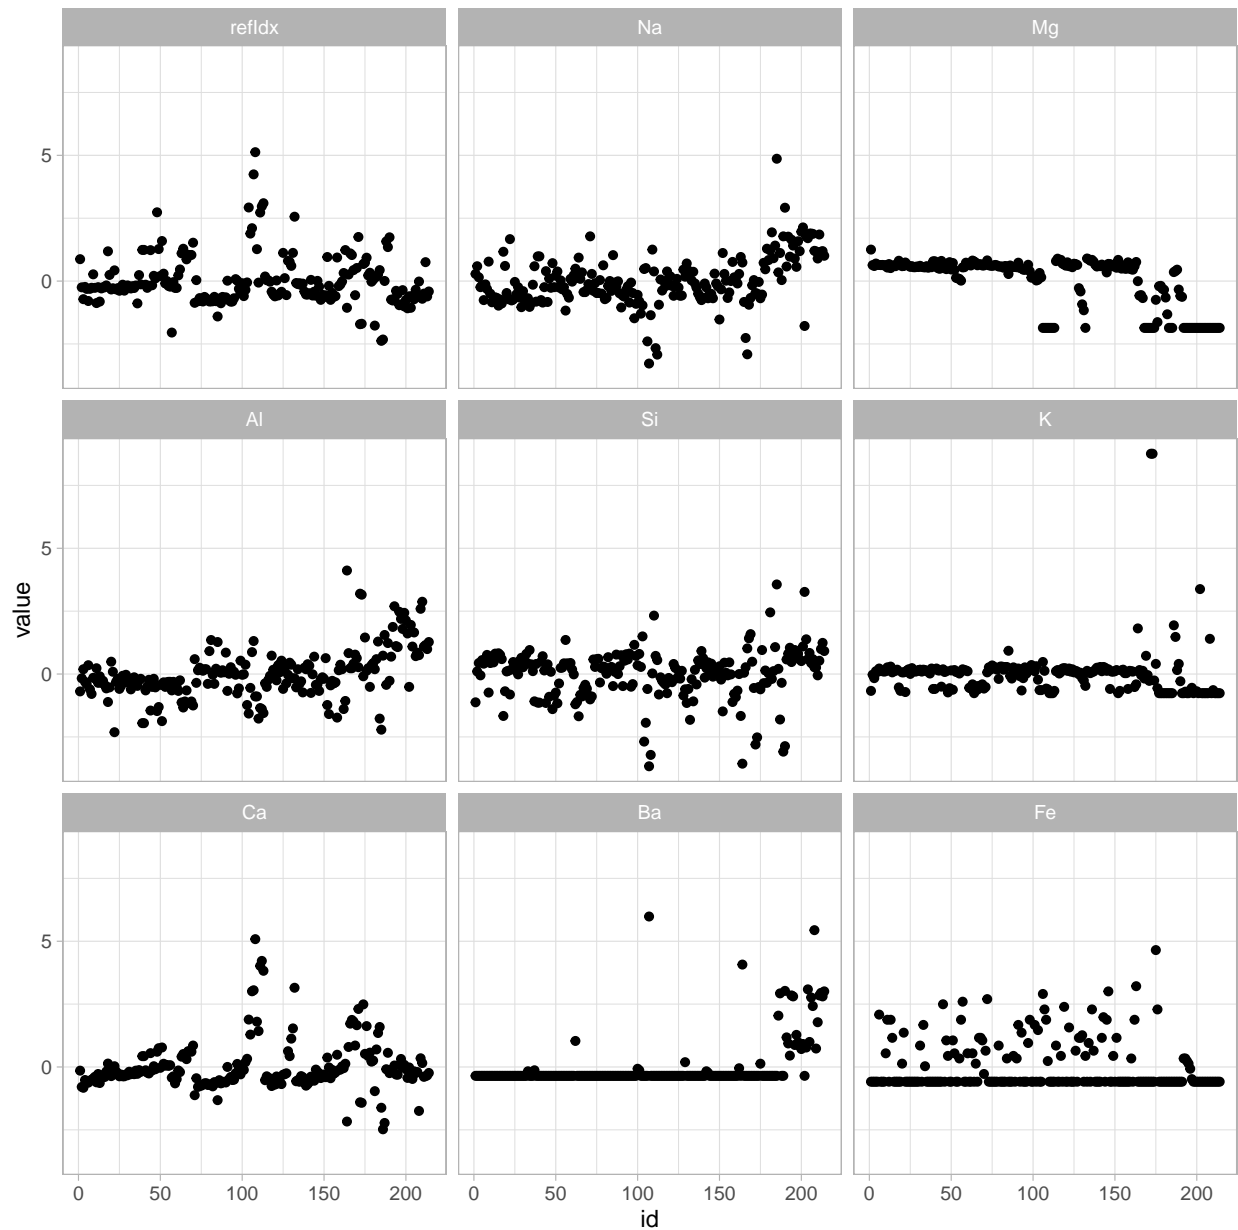
Based on the criteria above, as long as any of the attribute has a value that is 2 standard deviations from the mean, this id is defined as an outlier. The result returned 53 cases that include outliers, it is about 1/4 of the data, we shouldn't just removed them all.

Below I used some basic visualization (scatter plot, box plot) to see if they are helpful to visually define the outlier, then I used local outlier factor to find out outliers based on density following the instruction provided in Module 3's class material.

```
# box plot to see if we can remove the outliers from there
## reshape the dataset
glassDataCheck <- cbind(glassDataNorm,
                        id = glassData$id)
glassDataCheck <- melt(glassDataCheck,
                      id.vars = "id",
                      variable.name = "attribute")
## box plot
ggplot(glassDataCheck, aes(x = attribute, y = value)) +
  geom_boxplot() +
  theme_light()
```



```
# scatter plot to see if we can remove the outliers from there
ggplot(glassDataCheck, aes(x = id, y = value)) +
  geom_point() +
  facet_wrap(~attribute) +
  theme_light()
```



From the box plot, there is still a lot of outliers in each attribute, if we remove them all, it will still be a big portion of the data. The scatter plot doesn't provide a lot of information. I don't think it will be good to remove ids just based on one attribute.

Here I find the outliers using local outlier factor as it detects if a data point has a substantial lower density than their neighbors.

```
# calculate the local outlier factor of the matrix
outlier.scores <- lofactor(glassDataNorm,
                           k = 4)

# find out the top 10 outlier.scores and define them as outliers in the data set
outliersLOF <- order(outlier.scores,
                      decreasing = TRUE)[1:10]

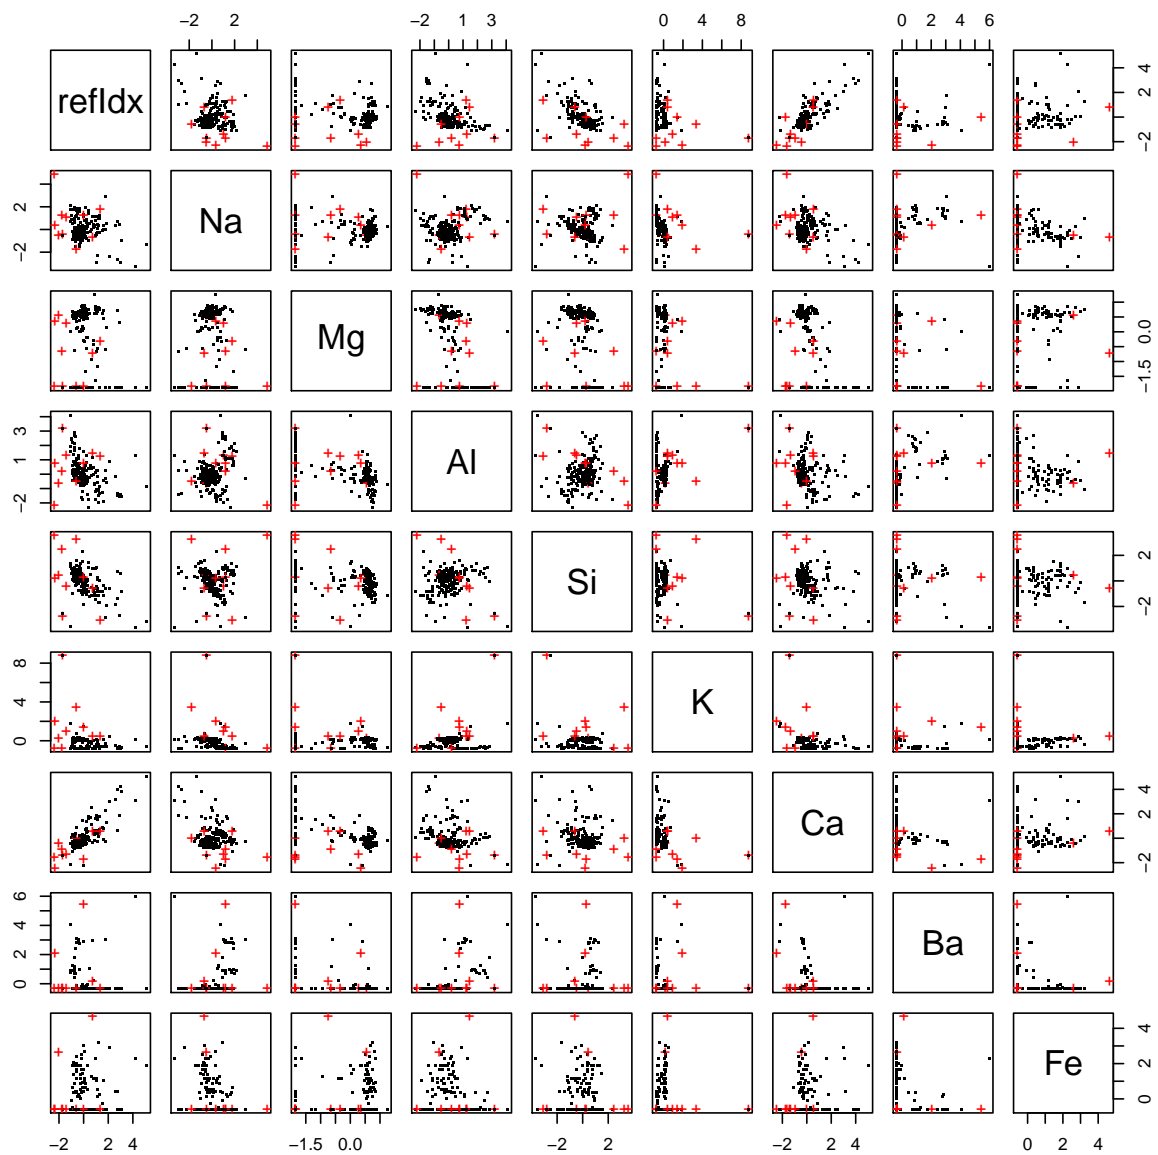
# plot the outliers with a pairs plot
n <- nrow(glassDataNorm)
```



```

labels <- 1:n
labels[-outliersLOF] <- "."
pch <- rep(".", n)
pch[outliersLOF] <- "+"
col <- rep("black", n)
col[outliersLOF] <- "red"
pairs(glassDataNorm,
      pch = pch,
      col = col)

```



Question 6

After removing the ID column (column 1), standardize the scales of the numeric columns, except the last one (the glass type), using z-score.

```
# the dataset has already been standardized above (in question 5),  
# add type column back and make sure it's normalized  
glassDataNorm$type <- glassData$type  
glassDataNorm$type <- as.factor(glassDataNorm$type)  
summary(glassDataNorm)
```

```
##      refIdx      Na      Mg      Al  
## Min.   :-2.3759 Min.   :-3.2793 Min.   :-1.8611 Min.   :-2.3132  
## 1st Qu.: -0.6069 1st Qu.: -0.6127 1st Qu.: -0.3948 1st Qu.: -0.5106  
## Median :-0.2257 Median :-0.1321 Median : 0.5515 Median :-0.1701  
## Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.0000  
## 3rd Qu.: 0.2608 3rd Qu.: 0.5108 3rd Qu.: 0.6347 3rd Qu.: 0.3707  
## Max.    : 5.1252 Max.    : 4.8642 Max.    : 1.2517 Max.    : 4.1162  
##      Si      K      Ca      Ba  
## Min.   :-3.6679 Min.   :-0.76213 Min.   :-2.4783 Min.   :-0.3521  
## 1st Qu.: -0.4789 1st Qu.: -0.57430 1st Qu.: -0.5038 1st Qu.: -0.3521  
## Median : 0.1795 Median : 0.08884 Median :-0.2508 Median :-0.3521  
## Mean   : 0.0000 Mean   : 0.00000 Mean   : 0.0000 Mean   : 0.0000  
## 3rd Qu.: 0.5636 3rd Qu.: 0.17318 3rd Qu.: 0.1515 3rd Qu.: -0.3521  
## Max.    : 3.5622 Max.    : 8.75961 Max.    : 5.0824 Max.    : 5.9832  
##      Fe      type  
## Min.   :-0.5851 1:70  
## 1st Qu.: -0.5851 2:76  
## Median :-0.5851 3:17  
## Mean   : 0.0000 5:13  
## 3rd Qu.: 0.4412 6: 9  
## Max.    : 4.6490 7:29
```

Question 7

The data set is sorted, so creating a validation data set requires random selection of elements. Create a stratified sample where you randomly select 15% of each of the cases for each glass type to be part of the validation data set. The remaining cases will form the training data set.

```
# randomly select 15% from each type  
glassDataValid <- glassDataNorm %>%  
  group_by(type) %>%  
  sample_frac(., 0.15)  
# get the remain data as training data  
glassDataTraining <- glassDataNorm %>%  
  anti_join(glassDataValid)
```

```
## Joining, by = c("refIdx", "Na", "Mg", "Al", "Si", "K", "Ca", "Ba", "Fe", "type")
```

Question 8

Implement the k-NN algorithm in R (do not use an implementation of k-NN from a package) and use your algorithm with a $k = 4$ to predict the glass type for the following two cases:

```

# create a data frame for the cases need to be predicted
glassDataTest_1 <- data.frame("id" = c(215, 216),
                             "refIdx" = c(1.51621, 1.57930),
                             "Na" = c(12.52, 12.69),
                             "Mg" = c(3.48, 1.86),
                             "Al" = c(1.39, 1.82),
                             "Si" = c(73.39, 72.62),
                             "K" = c(0.60, 0.52),
                             "Ca" = c(8.55, 10.52),
                             "Ba" = c(0.00, 0.00),
                             "Fe" = c(0.07, 0.05),
                             "type" = c(0, 0))

# normalize the dataset
glassDataTest_1 <- rbind(glassData,
                         glassDataTest_1)
glassDataTest_1[, 2:10] <- as.data.frame(scale(glassDataTest_1[,2:10]))
glassDataTest_1 <- glassDataTest_1 %>%
  filter(id == 215 | id == 216) %>%
  select(-id)
# duplicate one for comparing with the class package
glassDataTest_2 <- glassDataTest_1
glassDataTest_2

```

```

##      refIdx      Na      Mg      Al      Si      K      Ca
## 1 -0.4728742 -1.0783026 0.5533436 -0.1133190 0.95238059 0.15767490 -0.2902021
## 2 11.8167501 -0.8700927 -0.5733715 0.7508082 -0.04428423 0.03444459 1.0963553
##      Ba      Fe type
## 1 -0.3502235 0.13365304 0
## 2 -0.3502235 -0.07255451 0

```

Create a knn function named “knn_1” to predict these two cases glass type.

```

knn_1 <- function(trainData, testData, k){

  for(i in 1:nrow(testData)){
    for(j in 1:(ncol(trainData)-1)){
      # calculate distance between each value in test data and each value in train data
      d <- (trainData[, j] - testData[i, j]) ^ 2
      # get the index of top k nearest neighbors
      knnIdx <- head(sort(d, index.return = TRUE)$ix, k)
      # get the type of glass of these k nearest neighbors
      nn <- sort(table(trainData[knnIdx, 10]), TRUE)
      # add the predicted type to test data's type column
      testData[i, 10] <- as.numeric(names(nn)[1])
    }
  }
  testData
}

# use this function to predict these two cases type of glass
glassDataPred_1 <- knn_1(glassDataTraining,
                        glassDataTest_1,
                        4)

glassDataPred_1

```

```
##      refIdx      Na      Mg      Al      Si      K      Ca
## 1 -0.4728742 -1.0783026 0.5533436 -0.1133190 0.95238059 0.15767490 -0.2902021
## 2 11.8167501 -0.8700927 -0.5733715 0.7508082 -0.04428423 0.03444459 1.0963553
##      Ba      Fe type
## 1 -0.3502235 0.13365304 1
## 2 -0.3502235 -0.07255451 1
```

Question 9

Apply the knn function from the class package with $k = 4$ and redo the cases from Question(8). Compare your answers.

```
set.seed(123)
# run knn function from class package
glassDataPred_2 <- knn(train = glassDataTraining,
                       test = glassDataTest_2,
                       cl = glassDataTraining$type,
                       k = 4)
glassDataPred_2
```

```
## [1] 1 2
## Levels: 1 2 3 5 6 7
```

```
# add the predicted results to the data frame
glassDataTest_2$type <- glassDataPred_2
glassDataTest_2
```

```
##      refIdx      Na      Mg      Al      Si      K      Ca
## 1 -0.4728742 -1.0783026 0.5533436 -0.1133190 0.95238059 0.15767490 -0.2902021
## 2 11.8167501 -0.8700927 -0.5733715 0.7508082 -0.04428423 0.03444459 1.0963553
##      Ba      Fe type
## 1 -0.3502235 0.13365304 1
## 2 -0.3502235 -0.07255451 2
```

My function predicted both of the cases are glass type 1, the knn function from class package predicted them as 1 and 2.

Question 10

Using k-NN from the class package, create a plot of k (x-axis) from 2 to 8 versus accuracy (percentage of correct classifications) using ggplot.

```
# create a data frame called accurach with one column for all k values,
# and one column for each k's accuracy
accuracy <- data.frame("k" = c(2, 3, 4, 5, 6, 7, 8),
                       "accuracy" = c(0, 0, 0, 0, 0, 0, 0))
# create a loop to calculate accuracy for each k
for (k in 2:8){

  set.seed(123)
```

```


# get the predict type
glassDataPred_3 <- knn(train = glassDataTraining,
                       test = glassDataValid,
                       cl = glassDataTraining$type,
                       k)

# create a matrix with type in validation data and type predicted by knn
mtrix <- as.matrix(table(glassDataValid$type,
                        glassDataPred_3))

# calculate accuracy and add it to the accuracy data frame
accuracy[which(k == accuracy$k), 2] <- sum(diag(mtrix)) / nrow(glassDataValid)
}

# plot the k-accuracy table
ggplot(accuracy, aes(x = k, y = accuracy)) +
  geom_point() +
  geom_line()+
  scale_x_continuous(breaks = seq(2, 8, by = 1)) +
  labs(x = "k", y = "Accuracy") +
  theme_light()

```



Question 11

Download this (modified) version of the glass dataset containing missing values in column 4. Identify the missing values. Impute the missing values of this continuous numeric column using your regression version of kNN from Problem 2 below using the other columns are predictor features.

```

# read in the data file
glassDataModified <- read.csv("da5030.glass.data_with_missing_values.csv",
                             header = FALSE)

# give columns names
colnames(glassDataModified) <- c("id", "refIdx", "Na", "Mg", "Al",
                                "Si", "K", "Ca", "Ba", "Fe", "type")

# normalize the dataset
glassDataModifiedNorm <- as.data.frame(scale(glassDataModified[,2:11]))

# filter out cases that contain missing Mg values
glassDataModifiedImpute <- glassDataModifiedNorm %>%
  filter(is.na(Mg)) %>%
  select(-Mg)

# identify the target data
glassDataModifiedTarget <- glassDataModifiedNorm %>%
  filter(!is.na(Mg)) %>%
  select(Mg)

# identify the training data
glassDataModifiedTrain <- glassDataModifiedNorm %>%
  filter(!is.na(Mg)) %>%
  select(-Mg)

# copy the knn.reg function from below
knn.reg <- function(new_data, target_data, train_data, k){
  # create a predict table for results
  pred <- as.data.frame(rep(0, nrow(new_data)))
  colnames(pred) <- "pred"

  for(i in 1:nrow(new_data)){
    for(j in 1:ncol(new_data)){
      # calculate distance between each value in test data and each value in train data
      d <- (train_data[, j] - as.numeric(new_data[i, j])) ^ 2
      # get the index of top k nearest neighbors
      knnIdx <- head(sort(d, index.return = TRUE)$ix, k)
      # get the type of glass of these k nearest neighbors
      nn <- sort(table(target_data[knnIdx, 1]), TRUE)
      # calculate the weighted products
      weighted <- 4 * as.numeric(names(nn)[1]) +
        2 * as.numeric(names(nn)[2]) +
        sum(as.numeric(names(nn)[3:k]))
      # calculate the weighted average
      weightedAve <- weighted / (4 + 2 + (k - 2))
      # assign the weighted average to predict table
      pred[i, 1] <- weightedAve
    }
  }
  pred
}

# run the knn.reg function
set.seed(123)
knn.reg(glassDataModifiedImpute,
        glassDataModifiedTarget,

```

```
glassDataModifiedTrain,
k = 4)
```

```
##          pred
## 1  0.685387979
## 2  0.685387979
## 3  0.633846248
## 4  0.577861264
## 5 -0.601377992
## 6 -0.340114735
## 7 -0.005982135
## 8 -0.005982135
## 9 -0.005982135
```

Above is the predicted Mg values for the missing data.

Problem 2

Question 1

Investigate this data set of home prices in King County (USA)

```
houseData <- read.csv("kc_house_data.csv")
# check missing values
any(is.na(houseData))
```

```
## [1] FALSE
```

```
# check structure
str(houseData)
```

```
## 'data.frame':    21613 obs. of  21 variables:
## $ id             : num  7.13e+09 6.41e+09 5.63e+09 2.49e+09 1.95e+09 ...
## $ date           : chr   "20141013T000000" "20141209T000000" "20150225T000000" "20141209T000000" ...
## $ price          : num   221900 538000 180000 604000 510000 ...
## $ bedrooms       : int    3 3 2 4 3 4 3 3 3 3 ...
## $ bathrooms      : num    1 2.25 1 3 2 4.5 2.25 1.5 1 2.5 ...
## $ sqft_living    : int   1180 2570 770 1960 1680 5420 1715 1060 1780 1890 ...
## $ sqft_lot       : int   5650 7242 10000 5000 8080 101930 6819 9711 7470 6560 ...
## $ floors         : num    1 2 1 1 1 1 2 1 1 2 ...
## $ waterfront     : int    0 0 0 0 0 0 0 0 0 0 ...
## $ view           : int    0 0 0 0 0 0 0 0 0 0 ...
## $ condition      : int    3 3 3 5 3 3 3 3 3 3 ...
## $ grade          : int    7 7 6 7 8 11 7 7 7 7 ...
## $ sqft_above     : int   1180 2170 770 1050 1680 3890 1715 1060 1050 1890 ...
## $ sqft_basement : int    0 400 0 910 0 1530 0 0 730 0 ...
## $ yr_built       : int   1955 1951 1933 1965 1987 2001 1995 1963 1960 2003 ...
## $ yr_renovated   : int    0 1991 0 0 0 0 0 0 0 0 ...
## $ zipcode        : int   98178 98125 98028 98136 98074 98053 98003 98198 98146 98038 ...
## $ lat            : num    47.5 47.7 47.7 47.5 47.6 ...
```

```
## $ long      : num  -122 -122 -122 -122 -122 ...
## $ sqft_living15: int   1340 1690 2720 1360 1800 4760 2238 1650 1780 2390 ...
## $ sqft_lot15  : int   5650 7639 8062 5000 7503 101930 6819 9711 8113 7570 ...
```

Question 2

Save the price column in a separate vector/dataframe called `target_data`. Move all of the columns except the ID, data, price, yr_renovated, zipcode, lat, long, sqft_living12, and sqft_lot 15 columns into a new data frame called `train_data`.

```
# save the price column to target_data
target_data <- houseData %>% select(price)
# move required columns to train_data
train_data <- houseData %>% select(-c(id, date, price, yr_renovated,
                                     zipcode, lat, long, sqft_living15,
                                     sqft_lot15))

# check column info of train_data
sapply(train_data, class)
```

```
##      bedrooms      bathrooms      sqft_living      sqft_lot      floors
##      "integer"      "numeric"      "integer"      "integer"      "numeric"
##      waterfront      view      condition      grade      sqft_above
##      "integer"      "integer"      "integer"      "integer"      "integer"
##      sqft_basement      yr_built
##      "integer"      "integer"
```

Question 3

Normalize all of the columns (except the boolean columns `waterfront` and `view`) using min-max normalization

```
# create the min-max function
minMax <- function(x){
  return((x - min(x)) / (max(x) - min(x)))
}

# move waterfront and view column to the end of the data set
train_data_norm <- train_data %>%
  relocate(c(waterfront, view),
           .after = last_col())

# normalize the data set using min-max function
train_data_norm[, 1:10] <- as.data.frame(lapply(train_data_norm[, 1:10],
                                                minMax))
```

Question 4

Build a function called `knn.reg` that implements a regression version of k-NN that averages the prices of the k nearest neighbors using a weighted average where the weight is 4 for the closest neighbor, 2 for the second closest and 1 for the remaining neighbors (recall that a weighted average requires that you divide the sum product of the weight and values by the sum of the weights).


```

# create the knn.reg function
knn.reg <- function(new_data, target_data, train_data, k){
  # create a predict table for results
  pred <- as.data.frame(rep(0, nrow(new_data)))
  colnames(pred) <- "pred"

  for(i in 1:nrow(new_data)){
    for(j in 1:ncol(new_data)){
      # calculate distance between each value in test data and each value in train data
      d <- (train_data[, j] - as.numeric(new_data[i, j])) ^ 2
      # get the index of top k nearest neighbors
      knnIdx <- head(sort(d, index.return = TRUE)$ix, k)
      # get the type of glass of these k nearest neighbors
      nn <- sort(table(target_data[knnIdx, 1]), TRUE)
      # calculate the weighted products
      weighted <- 4 * as.numeric(names(nn)[1]) +
        2 * as.numeric(names(nn)[2]) +
        sum(as.numeric(names(nn)[3:k]))
      # calculate the weighted average
      weightedAve <- weighted / (4 + 2 + (k - 2))
      # assign the weighted average to predict table
      pred[i, 1] <- weightedAve
    }
  }
  pred
}

```

Question 5

Forecast the price of this new home using your regression k-NN using $k = 4$.

```

# type in the data
new_data <- data_frame("bedrooms" = 4,
  "bathrooms" = 3,
  "sqft_living" = 4852,
  "sqft_lot" = 11245,
  "floors" = 3,
  "condition" = 3,
  "grade" = 11,
  "sqft_above" = 2270,
  "sqft_basement" = 820,
  "yr_built" = 1986,
  "waterfront" = 1,
  "view" = 1)

```

```

## Warning: 'data_frame()' is deprecated as of tibble 1.1.0.
## Please use 'tibble()' instead.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_warnings()' to see where this warning was generated.

```

```

# normalize the new data
new_data <- rbind(train_data, new_data)

```

```

new_data <- new_data %>%
  relocate(c(waterfront, view),
            .after = last_col())
new_data[, 1:10] <- as.data.frame(lapply(new_data[, 1:10],
                                         minMax))

new_data <- new_data %>% slice(n())
# run the knn.reg function
set.seed(123)
knn.reg(new_data,
        target_data,
        train_data,
        4)

```

```

##          pred
## 1 584437.5

```

As the result from knn.reg, the predicted price of this new home is 584437.5